

Multi-Tenancy on Rails

An Ann Arbor Ruby Brigade Feature Length Presentation

you can go home when the number in the bottom right is 56ish

i know you cant see the numbers now but wait a few slides

John DeSilva

April 25

Revela, Inc.

What is Multi-Tennancy?

What is Multi-Tennancy?

The term “software multitenancy” refers to a software architecture in which a single instance of software runs on a server and serves multiple tenants. A tenant is a group of users who share a common access with specific privileges to the software instance.

- Wikipedia

What is Multi-Tennancy?

In short, we want to scale our customer base by sliding the heroku slider to the right instead of creating new heroku apps with a new db per customer.

if we need to add databases to scale to our data size or for customer reasons,

thats fine,

but only one app cluster

(or a multi zone federated cluster)

Types of Tennancy

'Multi User' Applications

Some applications are used by many individual users. Often B2C.

- Twitter
- Gmail
- Spotify

'Tenantless' Applications

Many applications do not have tenancy at all. Lots of low touch news, marketing, and cooking sites here.

- Allrecipes
- 4chan
- Darksky
- BuzzFeed
- Wikipedia

'Organization' Applications

Often, enterprise applications have a 'shared dataset' amongst users of different organizations. B2B SaaS products tend to fall into this bucket.

- Basecamp
- Slack
- Salesforce

'Gray Area' Applications

Some applications do not clearly fall into one of the previous three categories.


- Github
- Shopify

```
$ git checkout no_tennancy
```

(1/2): Query Scoping Tennancy

Straightforward

```
class RecipesController
  def show
    @recipe = Recipe.find(params[:id])
  end
end
```



Pretty Straightforward

```
class RecipesController
  def show
    @recipe = current_user.recipes
                               .find(params[:id])
  end
end
```

Kinda Straightforward

```
class RecipesController
  def show
    @recipe = current_user
               .company
               .recipes
               .find(params[:id])
  end
end
```

Somewhat Straightforward

```
class RecipesController
  before_action :set_recipe # Rails-it-up-a-notch!

  def show; end

  def set_recipe
    @recipe = recipes.find(params[:id])
  end

  delegate :company, to: :current_user
  delegate :recipes, to: :company
end
```



```
$ git checkout scoping
```



**MUCH
LATER**

Companies

\-> Buildings

 \-> Units

 \-> Leases

 \-> Lease Memberships

 \-> Tenants

 \-> Payments

‘Display a list of Payments’

```
select * from payments
join tenants on tenants.id = payments.tenant_id
join lms on lms.tenant_id = tenants.id
join leases on leases.id = lms.lease_id
join units on units.id = leases.unit_id
join buildings on buildings.id = units.building_id
join companies on companies.id = buildings.company_id
where companies.id = $1;
```

```
$1 => current_user.company_id;
```



y tho

Why not denormalize `customer_id`?

- The database schema should (probably?) represent your application's data, not your application's data performance needs or access patterns.
 - *This is CRUD, not Twitter*
- Now you have to have a Customer at arms reach at the point of Payment creation.
 - *You might, but hey - think of the children. Erm, I mean, background jobs and one off scripts*
- Hey, what else could use a `customer_id`?
 - *It will start to get sprinkled on everywhere, I promise*

Why not put `customer_id` on every table?

I'm pretty sure Citus Data's commercial `pg_shard` evolution actually does this or something similar

2



SCHEMA BASED MULTI TENNANCY

Schema

Database Schema

The term 'schema' refers to the organization of data as a blueprint of how the database is constructed (divided into database tables in the case of relational databases). The formal definition of a database schema is a set of formulas (sentences) called integrity constraints imposed on a database.

- Wikipedia

Lingo - Database Schema

```
# db/schema.rb
```

```
ActiveRecord::Schema.define(  
  version: 20170424000157  
) do
```

```
  create_table "companies", force: :cascade do |t|  
    t.string "name"
```

```
  # ...
```

Lingo - Database Schema

```
soda $ psql soda_development
```

```
[local] john@soda_development=# \dt
```

List of relations

Schema	Name	Type	Owner
public	ar_internal_metadata	table	john
public	companies	table	john
public	recipes	table	john
public	schema_migrations	table	john

(4 rows)

Lingo - Database Schema

```
[local] john@soda_development=# \d recipes
```

Table "public.recipes"

Column	Type	Modifi
id	integer	not n
title	character varying	
instructions	character varying	
created_at	timestamp without time zone	not n
updated_at	timestamp without time zone	not n
company_id	integer	

PostgreSQL Schema

A database contains one or more named schemas, which in turn contain tables. Schemas also contain other kinds of named objects, including data types, functions, and operators. The same object name can be used in different schemas without conflict; for example, both schema1 and myschema can contain tables named mytable. Unlike databases, schemas are not rigidly separated: a user can access objects in any of the schemas in the database they are connected to, if they have privileges to do so.

- <https://www.postgresql.org/docs/9.6/static/ddl-schemas.html>

There are several reasons why one might want to use schemas:

- To allow many users to use one database without interfering with each other.
- To organize database objects into logical groups to make them more manageable.
- Third-party applications can be put into separate schemas so they do not collide with the names of other objects.

```
[local] john@soda_development=# \dn
```

List of schemas

Name	Owner
public	postgres

(1 row)


```
[local] john@soda_development=# show search_path;
```

```
search_path
```

```
-----
```

```
"$user", public
```

```
(1 row)
```

```
Time: 0.114 ms
```

Mr. Krabs I Have An Idea



Lingo - PostgreSQL Schema

```
[local] john@soda_development=# \dn
```

List of schemas

Name		Owner
-----+		-----
coke		john
pepsi		john
public		postgres

(3 rows)

Lingo - PostgreSQL Schema

```
[local] john@soda_development=# \dt *.*
```

List of relations

Schema	Name	Type	Owner
-----+-----+-----+-----			
coke	recipes	table	john
pepsi	recipes	table	john
public	recipes	table	john

...

(74 rows)

```
[local] john@soda_development
```

```
# select title from recipes;
```

```
title
```

```
-----
```

```
(0 rows)
```

```
Time: 0.405 ms
```

Coke Recipes

```
[local] john@soda_development
```

```
# select title from coke.recipes;
```

```
      title
```

```
-----
```

```
Coca-Cola
```

```
Sprite
```

```
Mellow Yellow
```

```
(3 rows)
```

```
Time: 0.355 ms
```

Pepsi Recipes

```
[local] john@soda_development
```

```
# select title from pepsi.recipes;
```

```
title
```

```
-----
```

```
Pepsi
```

```
Mountain Dew
```

```
Mug Root Beer
```

```
(3 rows)
```

```
Time: 0.292 ms
```

Pepsi Recipes, Magically

```
[local] john@soda_development  
# set search path to pepsi;  
# select title from recipes;
```

title

Pepsi

Mountain Dew

Mug Root Beer

(3 rows)

Time: 0.292 ms

The apartment Gem

 [influitive/apartment](#)

Database multi-tenancy for Rack (and Rails) applications

```
Apartment::Tenant.create('pepsi')  
# => create schema pepsi;
```

```
Apartment::Tenant.switch!('pepsi')  
# => set search_path to pepsi;
```

```
Apartment::Tenant.drop('pepsi')  
# => drop schema pepsi cascade;
```

Installing apartment

```
# Gemfile
```

```
gem 'apartment'
```

```
# and then do a
```

```
rails g apartment:install
```

Configuring apartment - Elevator

```
# config/initializers/apartment.rb

require 'apartment/elevators/subdomain'

Rails.application.config.middleware.use \
  'Apartment::Elevators::Subdomain'

# (this is already the default config)
```

Configuring apartment - Subdomain Elevator

GET pepsi.recipevault.com/recipes/1

=> set search_path to 'pepsi';

Recipe.find(1)

=> select * from recipes
where recipes.id = 1;

(aka)

=> select * from pepsi.recipes
where recipes.id = 1;

Configuring apartment - Public Models

```
# config/initializers/apartment.rb
```

```
Apartment.configure do |config|
```

```
# Keep companies in the public schema
```

```
config.excluded_models = %w(Company)
```

```
end
```

Configuring apartment - Public Models

```
Apartment::Tenant.switch!('coke')
```

```
Company.all
```

```
=> select * from public.companies;  
      ^^^^^
```

```
Recipe.all
```

```
=> select * from recipes;
```

```
(aka)
```

```
=> select * from coke.recipes;  
      ^^^^^
```


Configuring apartment - Schemas

```
# config/initializers/apartment.rb
Apartment.configure do |config|

  # Find subdomains from Companies
  config.tenant_names = lambda do
    Company.pluck(:subdomain)
  end

end
```

`pepsi.localhost:3000/recipes`

Accessing Local Tenants - lvh.me

```
~ % dig +noall +answer lvh.me @8.8.8.8
```

```
lvh.me. 1544 IN A 127.0.0.1
```

Accessing Local Tenants

`pepsi.lvh.me:3000/recipes`

instead of

`localhost:3000/recipes`

```
$ git checkout apartment
```

Caveats

Migrations

```
# Every schema has to be migrated individually.  
# Apartment does this transparently,  
# but it could be problematic with  
# thousands of tenants.  
#  
# It's probably fine.
```

```
# It seems annoying at first  
Company.find_by(subdomain: 'pepsi').activate!  
# but then you realize its worth it
```


Background Jobs

```
# For normal jobs: #solved  
#  
# This adds metadata containing the current  
# tenant to the job payload, then transparently  
# switch!s upon running the job  
gem 'apartment-sidekiq'
```

Background Jobs

```
# For system jobs:
class MuhJob
  def perform(*args)
    Company.find_each do |company|
      company.activate!

      # Normal job stuff
      Post.old.archive!
    end
  end
end
```

Action Cable

It's a pretty quick fix.

```
class ApplicationCable::Connection < ActionCable::C
  identified_by :current_user, :tenant
  #                                     ^^^^^^^
```

```
def connect
  self.tenant = request.subdomain
  Apartment::Tenant.switch!(tenant)

  # As usual
  self.current_user = find_current_user
end
```

Considerations

Multi-Tennancy And You

Think about what kind of tennancy model your application will have to support. This seems obvious, but its actually so obvious that you might overlook it and mess up. (It took me three tries to get this right)

Different types of tennancy do not necessarily correspond to distinct architecture solutions. Although Salesforce and Basecamp are both B2B SaaS organizational products, they have very different use cases, business goals, and data requirements.

Thinking about user behaviour should tell you what kind of tennancy you should support, but your data complexity and business requirements should guide you when implementing it.